

Mechanistic Anatomy of Memory-Retrieval Failure as Context Fills

Kunj Rathod¹

¹University of Utah — MEMROT v2

June 2026

Abstract

Long-context language models lose access to stored facts well before their context window is exhausted (“context rot”), yet the underlying mechanism has been reported only correlationally. Using **Gemma-2-2B-it** within its native 8k window, we build *controlled haystacks* from LongMemEval-S material: the same question is evaluated at every fill level $k \in \{0, 1, 2, 4, 8, 14\}$ distractor sessions, with the gold session’s position randomized and logged. On 140 questions, accuracy falls monotonically from 0.71 at $k=0$ to 0.34 at $k=14$. Retrieval-head attention mass on the gold span and recall-associated sparse-autoencoder (SAE) feature activations both decline with k (mixed-effects $\beta_k = -0.023$, $p=2.4 \times 10^{-191}$; feature $\beta_k = -0.535$, $p=3.0 \times 10^{-90}$), controlling for gold position. Crucially, in a *paired within-question* design the fill level at which the internal signal collapses predicts the level at which the answer flips (Spearman $\rho = 0.731$, $p=3.6 \times 10^{-9}$, $n=48$; Cox HR = 5.3×10^{-4} , $p=3.6 \times 10^{-9}$). Causally, mean-ablating the 20 identified retrieval heads at $k=0$ drops accuracy by 0.82 (random-head control: 0.25), and feature-mode intervention at high k recovers 0.161 (95% CI [0.032, 0.290]) of 31 flipped answers. The pipeline is seeded and re-runnable byte-identically.

1 Introduction

Context rot — the degradation of retrieval as the context fills, well inside the model’s nominal window — has been documented behaviorally but not mechanistically. Two interpretability instruments let us watch the failure as it happens: *retrieval heads*, attention heads that copy from the evidence span into the generated answer, and *SAE features*, sparse directions that activate when the model recalls a fact. We ask three questions:

- H1** Does the internal retrieval machinery (attention mass on gold, recall-feature activation) degrade as context fills?
- H2** Does the *per-question* collapse point of that internal signal predict the *per-question* behavioral flip point?
- H3** Is the machinery causally responsible for the correct answer?

The novelty over prior correlational context-rot reports is a paired within-question design in which each question is its own control, isolating the effect of context fill from the effect of question difficulty.

2 Method: controlled-haystack paired design

Figure 1 summarizes the pipeline. From LongMemEval-S we keep the 140 questions that (i) are non-abstention, (ii) have *exactly one* gold evidence session of $\leq 2,000$ tokens (long gold sessions are trimmed to a turn window around the **has_answer** evidence turns), and (iii) are not of type **single-session-preference**, whose gold “answers” are prose preference-rubrics that the heuristic substring/F1 grader cannot score.

For each question and each k we pack the gold session plus k topic-matched distractor sessions (each truncated to a fixed per-distractor cap so context grows roughly linearly with k) into $\leq 7,500$ tokens, placing the gold session at a uniformly random, logged position. A question enters the H2/H3 analysis only if it is answered correctly at $k=0$, so any failure at $k > 0$ is attributable to context fill, not missing knowledge.

We instrument Gemma-2-2B-it with eager attention. Retrieval heads are identified by copy-score probes built from benchmark vocabulary (20 heads selected); the five SAE layers {20, 21, 22, 23, 24} are chosen by an empirical separability scan over all 26 Gemma Scope **res** layers. Decoding is greedy; seed 0; CUBLAS_WORKSPACE_CONFIG=:4096:8 with deterministic algorithms makes a re-executed run byte-identical

Table 1: Accuracy by fill level k on 140 questions. Accuracy decays monotonically as distractor sessions are added, despite the gold evidence remaining present and in-window.

k	0	1	2	4	8	14
Acc.	0.714	0.671	0.579	0.586	0.521	0.343

on `answers.parquet`.

3 Results

H1 (descriptive). Retrieval-head attention mass on the gold span declines with k ($\beta_k = -0.023$, $p = 2.4 \times 10^{-191}$; mixed-effects model with question as a random intercept and gold position as a covariate, gold-position $\beta = 0.446$). Recall-feature activation shows the same decline ($\beta_k = -0.535$, $p = 3.0 \times 10^{-90}$). See Figures 2 and 5.

H2 (paired, headline). Among the 100 questions correct at $k = 0$, the per-question signal-collapse level predicts the answer-flip level: Spearman $\rho = 0.731$ ($p = 3.6 \times 10^{-9}$, $n = 48$ questions with both events observed). A survival model over all questions (with censoring) gives a hazard ratio of 5.3×10^{-4} ($p = 3.6 \times 10^{-9}$) per unit normalized signal-AUC. The feature-signal variant gives $\rho = 0.763$. See Figure 3.

H3 (causal). Mean-ablating the 20 retrieval heads at $k = 0$ on the 100 previously-correct questions drops accuracy by 0.82, versus 0.25 for matched random heads — a $3.3\times$ specific effect. Intervening at high k in feature mode (SAE recall-feature re-injection) recovers 0.161 (95% CI [0.032, 0.290]) of the 31 flipped answers. See Figure 4.

4 Limitations

A 2B base model and a single benchmark source; eager-attention requirement (no flash-attention numerics); heuristic (non-LLM-judge) grading; and distractor construction choices (fixed per-distractor token cap; pool drawn from the benchmark’s own topic-matched filler). Gemma-2 interleaves sliding-window and global attention layers, so gold spans beyond the 4,096-token sliding window are structurally invisible to sliding layers; we log this but do not model it separately.

5 Related work

Context rot (Chroma); lost-in-the-middle (Liu et al.); retrieval heads (Wu et al.); Gemma Scope SAEs (Lieberum et al.); LongMemEval (Wu et al.).

6 Reproducibility

Seed 0; artifact schema v2; environment hash `d6e719d3`. The producing run is git `54907d9` plus the smoke-gate config/ data-prep amendment now committed as `83a762c` (model $\rightarrow$ `gemma-2-2b-it`; exclude `single-session-preference`; $n=140$). Three SLURM jobs reproduce everything end-to-end; Run 1 re-execution is byte-identical on `answers.parquet`.




Figure 1: Controlled-haystack paired pipeline. Each question is replayed at six fill levels with the gold session at a randomized position; an instrumented forward pass records both the internal retrieval signal and the graded answer, which are aligned per-question to test whether internal collapse predicts the behavioral flip (H2) and whether the machinery is causal (H3).




Figure 2: H1. Internal retrieval signal (attention mass on the gold span) versus fill level k , with question as random intercept. The signal decays as context fills even though the gold evidence is present and in-window.




Figure 3: H2 (headline). Per-question internal-collapse level versus behavioral flip level. The two are tightly coupled ($\rho = 0.731$), supporting a mechanistic rather than incidental account of context rot.




Figure 4: H3 (causal). Left: accuracy drop from mean-ablating the 20 retrieval heads (0.82) versus matched random heads (0.25). Right: fraction of flipped answers recovered by feature-mode intervention (0.161, 95% CI [0.032, 0.290]).




Figure 5: H1 (features). Recall-associated SAE feature activation across layers {20–24} as a function of k ; activation weakens with fill, mirroring the attention-mass decline ($\beta_k = -0.535$).